

```

        printf("%s destructed.\n", name); ❷
    }
private:
    const char* const name;
};

void consumer(SimpleUniquePointer<Tracer> consumer_ptr) {
    printf("(cons) consumer_ptr: 0x%p\n", consumer_ptr.get()); ❸
}

int main() {
    auto ptr_a = SimpleUniquePointer(new Tracer{ "ptr_a" });
    printf("(main) ptr_a: 0x%p\n", ptr_a.get()); ❹
    consumer(std::move(ptr_a));
    printf("(main) ptr_a: 0x%p\n", ptr_a.get()); ❺
}
-----
ptr_a constructed. ❶
(main) ptr_a: 0x000001936B5A2970 ❹
(cons) consumer_ptr: 0x000001936B5A2970 ❸
ptr_a destructed. ❷
(main) ptr_a: 0x0000000000000000 ❺

```

Listing 6-15: A program investigating SimpleUniquePointers with the Tracer class

First, you dynamically allocate a Tracer with the message `ptr_a`. This prints the first message ❶. You use the resulting Tracer pointer to construct a `SimpleUniquePointer` called `ptr_a`. Next, you use the `get()` method of `ptr_a` to retrieve the address of its Tracer, which you print ❹. Then you use `std::move` to relinquish the Tracer of `ptr_a` to the `consumer` function, which moves `ptr_a` into the `consumer_ptr` argument.

Now, `consumer_ptr` owns the Tracer. You use the `get()` method of `consumer_ptr` to retrieve the address of Tracer, then print ❸. Notice this address matches the one printed at ❹. When `consumer` returns, `consumer_ptr` dies because its storage duration is the scope of `consumer`. As a result, `ptr_a` gets destructed ❷.

Recall that `ptr_a` is in a moved-from state—you moved its Tracer into `consumer`. You use the `get()` method of `ptr_a` to illustrate that it now holds a `nullptr` ❺.

Thanks to `SimpleUniquePointer`, you won't leak a dynamically allocated object; also, because the `SimpleUniquePointer` is just carrying around a pointer under the hood, move semantics are efficient.

NOTE *The `SimpleUniquePointer` is a pedagogical implementation of the `stdlib`'s `std::unique_ptr`, which is a member of the family of RAII templates called smart pointers. You'll learn about these in Part II.*

Type Checking in Templates

Templates are type safe. During template instantiation, the compiler pastes in the template parameters. If the resulting code is incorrect, the compiler will not generate the instantiation.

Consider the template function in Listing 6-16, which squares an element and returns the result.

```
template<typename T>
T square(T value) {
    return value * value; ❶
}
```

Listing 6-16: A template function that squares a value

The `T` has a silent requirement: it must support multiplication ❶.

If you try to use `square` with, say, a `char*`, the compilation will fail, as shown in Listing 6-17.

```
template<typename T>
T square(T value) {
    return value * value;
}

int main() {
    char my_char{ 'Q' };
    auto result = square(&my_char); ❶ // Bang!
}
```

Listing 6-17: A program with a failed template instantiation. (This program fails to compile.)

Pointers don't support multiplication, so template initialization fails ❶.

The `square` function is trivially small, but the failed template initialization's error message isn't. On MSVC v141, you get this:

```
main.cpp(3): error C2296: '*': illegal, left operand has type 'char *'
main.cpp(8): note: see reference to function template instantiation 'T
*square<char*>(T)' being compiled
        with
        [
            T=char *
        ]
main.cpp(3): error C2297: '*': illegal, right operand has type 'char *'
```

And on GCC 7.3, you get this:

```
main.cpp: In instantiation of 'T square(T) [with T = char*]':
main.cpp:8:32:   required from here
main.cpp:3:16: error: invalid operands of types 'char*' and 'char*' to binary
'operator*'
    return value * value;
           ~~~~~^~~~~~
```

These error messages exemplify the notoriously cryptic error messages emitted by template initialization failures.

Although template instantiation ensures type safety, the checking happens very late in the compilation process. When the compiler instantiates

a template, it pastes the template parameter types into the template. After type insertion, the compiler attempts to compile the result. If instantiation fails, the compiler emits the dying words inside the template instantiation.

C++ template programming shares similarities with *duck-typed languages*. Duck-typed languages (like Python) defer type checking until runtime. The underlying philosophy is that if an object looks like a duck and quacks like a duck, then it must be type duck. Unfortunately, this means you can't generally know whether an object supports a particular operation until you execute the program.

With templates, you cannot know whether an instantiation will succeed until you try to compile it. Although duck-typed languages might blow up at runtime, templates might blow up at compile time.

This situation is widely regarded as unacceptable by right-thinking people in the C++ community, so there is a splendid solution called concepts.

Concepts

Concepts constrain template parameters, allowing for parameter checking at the point of instantiation rather than the point of first use. By catching usage issues at the point of instantiation, the compiler can give you a friendly, informative error code—for example, “You tried to instantiate this template with a `char*`, but this template requires a type that supports multiplication.”

Concepts allow you to express requirements on template parameters directly in the language.

Unfortunately, concepts aren't yet officially part of the C++ standard, although they've been voted into C++ 20. At press time, GCC 6.0 and later support the Concepts Technical Specification, and Microsoft is actively working toward implementing concepts in its C++ compiler, MSVC. Regardless of its unofficial status, it's worth exploring concepts in some detail for a few reasons:

- They'll fundamentally change the way you achieve compile-time polymorphism. Familiarity with concepts will pay major dividends.
- They provide a conceptual framework for understanding some of the makeshift solutions that you can put in place to get better compiler errors when templates are misused.
- They provide an excellent conceptual bridge from compile-time templates to interfaces, the primary mechanism for runtime polymorphism (covered in Chapter 5).
- If you can use GCC 6.0 or later, concepts *are* available by turning on the `-fconcepts` compiler flag.

WARNING

C++ 20's final concept specification will almost certainly deviate from the Concepts Technical Specification. This section presents concepts as specified in the Concepts Technical Specification so you can follow along.

Defining a Concept

A concept is a template. It's a constant expression involving template arguments, evaluated at compile time. Think of a concept as one big *predicate*: a function that evaluates to true or false.

If a set of template parameters meets the criteria for a given concept, that concept evaluates to true when instantiated with those parameters; otherwise, it will evaluate to false. When a concept evaluates to false, template instantiation fails.

You declare concepts using the keyword `concept` on an otherwise familiar template function definition:

```
template<typename T1, typename T2, ...>
concept bool ConceptName() {
    --snip--
}
```

Type Traits

Concepts validate type parameters. Within concepts, you manipulate types to inspect their properties. You can hand roll these manipulations, or you can use the type support library built into the `stdlib`. The library contains utilities for inspecting type properties. These utilities are collectively called *type traits*. They're available in the `<type_traits>` header and are part of the `std` namespace. Table 6-1 lists some commonly used type traits.

NOTE

See Chapter 5.4 of *The C++ Standard Library, 2nd Edition*, by Nicolai M. Josuttis for an exhaustive listing of type traits available in the `stdlib`.

Table 6-1: Selected Type Traits from the `<type_traits>` Header

Type trait	Checks if template argument is . . .
<code>is_void</code>	<code>void</code>
<code>is_null_pointer</code>	<code>nullptr</code>
<code>is_integral</code>	<code>bool</code> , a <code>char</code> type, an <code>int</code> type, a <code>short</code> type, a <code>long</code> type, or a <code>long long</code> type
<code>is_floating_point</code>	<code>float</code> , <code>double</code> , or <code>long double</code>
<code>is_fundamental</code>	Any of <code>is_void</code> , <code>is_null_pointer</code> , <code>is_integral</code> , or <code>is_floating_point</code>
<code>is_array</code>	An array; that is, a type containing square brackets <code>[]</code>
<code>is_enum</code>	An enumeration type (<code>enum</code>)
<code>is_class</code>	A class type (but not a union type)
<code>is_function</code>	A function
<code>is_pointer</code>	A pointer; function pointers count, but pointers to class members and <code>nullptr</code> do not
<code>is_reference</code>	A reference (either lvalue or rvalue)
<code>is_arithmetic</code>	<code>is_floating_point</code> or <code>is_integral</code>

Type trait	Checks if template argument is . . .
<code>is_pod</code>	A plain-old-data type; that is, a type that can be represented as a data type in plain C
<code>is_default_constructible</code>	Default constructible; that is, it can be constructed without arguments or initialization values
<code>is_constructible</code>	Constructible with the given template parameters: this type trait allows the user to provide additional template parameters beyond the type under consideration
<code>is_copy_constructible</code>	Copy constructible
<code>is_move_constructible</code>	Move constructible
<code>is_destructible</code>	Destructible
<code>is_same</code>	The same type as the additional template parameter type (including <code>const</code> and <code>volatile</code> modifiers)
<code>is_invocable</code>	Invocable with the given template parameters: this type trait allows the user to provide additional template parameters beyond the type under consideration

Each type trait is a template class that takes a single template parameter, the type you want to inspect. You extract the results using the template's static member value. This member equals `true` if the type parameter meets the criteria; otherwise, it's `false`.

Consider the type trait classes `is_integral` and `is_floating_point`. These are useful for checking if a type is (you guessed it) integral or floating point. Both of these templates take a single template parameter. The example in Listing 6-18 investigates type traits with several types.

```
#include <type_traits>
#include <cstdio>
#include <cstdint>

constexpr const char* as_str(bool x) { return x ? "True" : "False"; } ❶

int main() {
    printf("%s\n", as_str(std::is_integral<int>::value)); ❷
    printf("%s\n", as_str(std::is_integral<const int>::value)); ❸
    printf("%s\n", as_str(std::is_integral<char>::value)); ❹
    printf("%s\n", as_str(std::is_integral<uint64_t>::value)); ❺
    printf("%s\n", as_str(std::is_integral<int&>::value)); ❻
    printf("%s\n", as_str(std::is_integral<int*>::value)); ❼
    printf("%s\n", as_str(std::is_integral<float>::value)); ❽
}

-----
True ❷
True ❸
True ❹
True ❺
False ❻
False ❼
False ❽
```

Listing 6-18: A program using type traits

Listing 6-18 defines the convenience function `as_str` ❶ to print Boolean values with the string `True` or `False`. Within `main`, you print the result of various type trait instantiations. The template parameters `int` ❷, `const int` ❸, `char` ❹, and `uint64_t` ❺ all return `true` when passed to `is_integral`. Reference types ❻ ❼ and floating-point types ❸ return `false`.

NOTE

Recall that `printf` doesn't have a format specifier for `bool`. Rather than using the integer format specifier `%d` as a stand-in, Listing 6-18 employs the `as_str` function, which returns the string literal `True` or `False` depending on the value of the `bool`. Because these values are string literals, you can capitalize them however you like.

Type traits are often the building blocks for a concept, but sometimes you need more flexibility. Type traits tell you *what* types are, but sometimes you must also specify *how* the template will use them. For this, you use requirements.

Requirements

Requirements are ad hoc constraints on template parameters. Each concept can specify any number of requirements on its template parameters. Requirements are encoded into `requires` expressions denoted by the `requires` keyword followed by function arguments and a body.

A sequence of syntactic requirements comprises the requirements expression's body. Each syntactic requirement puts a constraint on the template parameters. Together, `requires` expressions have the following form:

```
requires (arg-1, arg-2, ...❶) {  
    { expression1❷ } -> return-type1❸;  
    { expression2 } -> return-type2;  
    --snip--  
}
```

`Requires` expressions take arguments that you place after the `requires` keyword ❶. These arguments have types derived from template parameters. The syntactic requirements follow, each denoted with `{ } ->`. You put an arbitrary expression within each of the braces ❷. This expression can involve any number of the arguments to the argument expression.

If an instantiation causes a syntactic expression not to compile, that syntactic requirement fails. Supposing the expression evaluates without error, the next check is whether the return type of that expression matches the type given after the arrow `->` ❸. If the expression result's evaluated type can't implicitly convert to the return type ❸, the syntactic requirement fails.

If any of the syntactic requirements fail, the `requires` expression evaluates to `false`. If all of the syntactic requirements pass, the `requires` expression evaluates to `true`.

Suppose you have two types, `T` and `U`, and you want to know whether you can compare objects of these types using the equality `==` and inequality `!=` operators. One way to encode this requirement is to use the following expression.

```
// T, U are types
requires (T t, U u) {
    { t == u } -> bool; // syntactic requirement 1
    { u == t } -> bool; // syntactic requirement 2
    { t != u } -> bool; // syntactic requirement 3
    { u != t } -> bool; // syntactic requirement 4
}
```

The `requires` expression takes two arguments, one each of types `T` and `U`. Each of the syntactic requirements contained in the `requires` expression is an expression using `t` and `u` with either `==` or `!=`. All four syntactic requirements enforce a `bool` result. Any two types that satisfy this `requires` expression are guaranteed to support comparison with `==` and `!=`.

Building Concepts from Requires Expressions

Because `requires` expressions are evaluated at compile time, concepts can contain any number of them. Try to construct a concept that guards against the misuse of `mean`. Listing 6-19 annotates some of the implicit requirements used earlier in Listing 6-10.

```
template<typename T>
T mean(T* values, size_t length) {
    T result{}; ❶
    for(size_t i{}; i<length; i++) {
        result ❷+= values[i];
    }
    ❸return result / length;
}
```

Listing 6-19: A relisting of 6-10 with annotations for some implicit requirements on `T`

You can see three requirements implied by this code:

- `T` must be default constructible ❶.
- `T` supports operator+= ❷.
- Dividing a `T` by a `size_t` yields a `T` ❸.

From these requirements, you could create a concept called `Averageable`, as demonstrated in Listing 6-20.

```
template<typename T>
concept bool Averageable() {
    return std::is_default_constructible<T>::value ❶
        && requires (T a, T b) {
            { a += b } -> T; ❷
            { a / size_t{ 1 } } -> T; ❸
        };
}
```

Listing 6-20: An `Averageable` concept. Annotations are consistent with the requirements and the body of `mean`.

You use the type trait `is_default_constructible` to ensure that `T` is default constructible ❶, that you can add two `T` types ❷, and that you can divide a `T` by a `size_t` ❸ and get a result of type `T`.

Recall that concepts are just predicates; you're building a Boolean expression that evaluates to true when the template parameters are supported and false when they're not. The concept is composed of three Boolean expressions AND-ed (&&) together: two type traits ❶❸ and a requires expression. If any of the three returns false, the concept's constraints are not met.

Using Concepts

Declaring concepts is a lot more work than using them. To use a concept, just use the concept's name in place of the `typename` keyword.

For example, you can refactor Listing 6-13 with the `Averageable` concept, as shown in Listing 6-21.

```
#include <cstdint>
#include <type_traits>

template<typename T>
concept bool Averageable() { ❶
    --snip--
}

template<Averageable❷ T>
T mean(const T* values, size_t length) {
    --snip--
}

int main() {
    const double nums_d[] { 1.0f, 2.0f, 3.0f, 4.0f };
    const auto result1 = mean(nums_d, 4);
    printf("double: %f\n", result1);

    const float nums_f[] { 1.0, 2.0, 3.0, 4.0 };
    const auto result2 = mean(nums_f, 4);
    printf("float: %f\n", result2);

    const size_t nums_c[] { 1, 2, 3, 4 };
    const auto result3 = mean(nums_c, 4);
    printf("size_t: %d\n", result3);
}

-----
double: 2.500000
float: 2.500000
size_t: 2
```

Listing 6-21: A refactor of Listing 6-13 using `Averageable`

After defining `Averageable` ❶, you just use it in place of `typename` ❷. No further modification is necessary. The code generated from compiling Listing 6-13 is identical to the code generated from compiling Listing 6-21.

The payoff is when you get to try to use `mean` with a type that is not `Averageable`: you get a compiler error at the point of instantiation. This produces much better compiler error messages than you would obtain from a raw template.

Look at the instantiation of `mean` in Listing 6-22 where you “accidentally” try to average an array of double pointers.

```
--snip--
int main() {
    auto value1 = 0.0;
    auto value2 = 1.0;
    const double* values[] { &value1, &value2 };
    mean(values❶, 2);
}
```

Listing 6-22: A bad template instantiation using a non-`Averageable` argument

There are several problems with using values ❶. What can the compiler tell you about those problems?

Without concepts, GCC 6.3 produces the error message shown in Listing 6-23.

```
<source>: In instantiation of 'T mean(const T*, size_t) [with T = const
double*; size_t = long unsigned int]':
<source>:17:17:   required from here
<source>:8:12: error: invalid operands of types 'const double*' and 'const
double*' to binary 'operator+'
    result += values[i]; ❶
    ~~~~~^~~~~~
<source>:8:12: error:   in evaluation of 'operator+=(const double*, const
double*)'
<source>:10:17: error: invalid operands of types 'const double*' and 'size_t'
{aka 'long unsigned int'} to binary 'operator/'
    return result / length; ❷
    ~~~~~^~~~~~
```

Listing 6-23: Error message from GCC 6.3 when compiling Listing 6-22

You might expect a casual user of `mean` to be extremely confused by this error message. What is `i` ❶? Why is a `const double*` involved in division ❷?

Concepts provide a far more illuminating error message, as Listing 6-24 demonstrates.

```
<source>: In function 'int main()':
<source>:28:17: error: cannot call function 'T mean(const T*, size_t) [with T
= const double*; size_t = long unsigned int]'
    mean(values, 2); ❶
    ^
<source>:16:3: note:   constraints not satisfied
    T mean(const T* values, size_t length) {
    ^~~~
<source>:6:14: note: within 'template<class T> concept bool Averageable()
[with T = const double*]'
```

```

concept bool Averageable() {
    ~~~~~
<source>:6:14: note:      with 'const double* a'
<source>:6:14: note:      with 'const double* b'
<source>:6:14: note: the required expression '(a + b)' would be ill-formed ❷
<source>:6:14: note: the required expression '(a / b)' would be ill-formed ❸

```

Listing 6-24: Error message from GCC 7.2 when compiling Listing 6-22 with concepts enabled

This error message is fantastic. The compiler tells you which argument (values) didn't meet a constraint ❶. Then it tells you that values is not Averageable because it doesn't satisfy two required expressions ❷ ❸. You know immediately how to modify your arguments to make this template instantiation successful.

When concepts incorporate into the C++ standard, it's likely that the stdlib will include many concepts. The design goal of concepts is that a programmer shouldn't have to define very many concepts on their own; rather, they should be able to combine concepts and ad hoc requirements within a template prefix. Table 6-2 provides a partial listing of some concepts you might expect to be included; these are borrowed from Andrew Sutton's implementation of concepts in the Origins Library.

NOTE

See <https://github.com/asutton/origin/> for more information on the Origins Library. To compile the examples that follow, you can install Origins and use GCC version 6.0 or later with the `-fconcepts` flag.

Table 6-2: The Concepts Contained in the Origins Library

Concept	A type that . . .
Conditional	Can be explicitly converted to bool
Boolean	Is Conditional and supports <code>!</code> , <code>&&</code> , and <code> </code> Boolean operations
Equality_comparable	Supports <code>==</code> and <code>!=</code> operations returning a Boolean
Destructible	Can be destroyed (compare <code>is_destructible</code>)
Default_constructible	Is default constructible (compare <code>is_default_constructible</code>)
Movable	Supports move semantics: it must be move assignable and move constructible (compare <code>is_move_assignable</code> , <code>is_move_constructible</code>)
Copyable	Supports copy semantics: it must be copy assignable and copy constructible (compare <code>is_copy_assignable</code> , <code>is_copy_constructible</code>)
Regular	Is default constructible, copyable, and Equality_comparable
Ordered	Is Regular and is totally ordered (essentially, it can be sorted)
Number	Is Ordered and supports math operations like <code>+</code> , <code>-</code> , <code>/</code> , and <code>*</code>
Function	Supports invocation; that is, you can call it (compare <code>is_invocable</code>)
Predicate	Is a Function and returns bool
Range	Can be iterated over in a range-based for loop

There are several ways to build constraints into a template prefix. If a template parameter is only used to declare the type of a function parameter, you can omit the template prefix entirely:

```
return-type function-name(Concept1❶ arg-1, ...) {  
    --snip--  
}
```

Because you use a concept rather than a typename to define an argument's type ❶, the compiler knows that the associated function is a template. You are even free to mix concepts and concrete types in the argument list. In other words, whenever you use a concept as part of a function definition, that function becomes a template.

The template function in Listing 6-25 takes an array of `Ordered` elements and finds the minimum.

```
#include <origin/core/concepts.hpp>  
size_t index_of_minimum(Ordered❶* x, size_t length) {  
    size_t min_index{};  
    for(size_t i{ 1 }; i<length; i++) {  
        if(x[i] < x[min_index]) min_index = i;  
    }  
    return min_index;  
}
```

Listing 6-25: A template function using the `Ordered` concept

Even though there's no template prefix, `index_of_minimum` is a template because `Ordered` ❶ is a concept. This template can be instantiated in the same way as any other template function, as demonstrated in Listing 6-26.

```
#include <cstdio>  
#include <cstdint>  
#include <origin/core/concepts.hpp>  
  
struct Goblin{};  
  
size_t index_of_minimum(Ordered* x, size_t length) {  
    --snip--  
}  
  
int main() {  
    int x1[] { -20, 0, 100, 400, -21, 5123 };  
    printf("%zu\n", index_of_minimum(x1, 6)); ❶  
  
    unsigned short x2[] { 42, 51, 900, 400 };  
    printf("%zu\n", index_of_minimum(x2, 4)); ❷  
  
    Goblin x3[] { Goblin{}, Goblin{} };  
    //index_of_minimum(x3, 2); ❸ // Bang! Goblin is not Ordered.  
}
```

4 ❶
0 ❷

Listing 6-26: A listing employing `index_of_minimum` from Listing 6-25. Uncommenting ❸ causes compilation to fail.

The instantiations for `int` ❶ and `unsigned short` ❷ arrays succeed because types are `Ordered` (see Table 6-2).

However, the `Goblin` class is not `Ordered`, and template instantiation would fail if you tried to compile ❸. Crucially, the error message would be informative:

```
error: cannot call function 'size_t index_of_minimum(auto:1*, size_t) [with auto:1 = Goblin; size_t = long unsigned int]'
      index_of_minimum(x3, 2); // Bang! Goblin is not Ordered.
                        ^
note:   constraints not satisfied
      size_t index_of_minimum(Ordered* x, size_t length) {
      ~~~~~
note: within 'template<class T> concept bool origin::Ordered() [with T =
Goblin]'
      Ordered()
```

You know that the `index_of_minimum` instantiation failed and that the issue is with the `Ordered` concept.

Ad Hoc Requires Expressions

Concepts are fairly heavyweight mechanisms for enforcing type safety. Sometimes, you just want to enforce some requirement directly in the template prefix. You can embed `requires` expressions directly into the template definition to accomplish this. Consider the `get_copy` function in Listing 6-27 that takes a pointer and safely returns a copy of the pointed-to object.

```
#include <stdexcept>

template<typename T>
    requires ❶ is_copy_constructible<T>::value ❷
T get_copy(T* pointer) {
    if (!pointer) throw std::runtime_error{ "Null-pointer dereference" };
    return *pointer;
}
```

Listing 6-27: A template function with an ad hoc `requires` expression

The template prefix contains the `requires` keyword ❶, which begins the `requires` expression. In this case, the type trait `is_copy_constructible` ensures that `T` is copyable ❷. This way, if a user accidentally tries to `get_copy` with a pointer that points to an uncopyable object, they'll be presented with a clear explanation of why template instantiation failed. Consider the example in Listing 6-28.

```

#include <stdexcept>
#include <type_traits>

template<typename T>
    requires std::is_copy_constructible<T>::value
T get_copy(T* pointer) { ❶
    --snip--
}

struct Highlander {
    Highlander() = default; ❷
    Highlander(const Highlander&) = delete; ❸
};

int main() {
    Highlander connor; ❹
    auto connor_ptr = &connor; ❺
    auto connor_copy = get_copy(connor_ptr); ❻
}

```

```

In function 'int main()':
error: cannot call function 'T get_copy(T*) [with T = Highlander]'
    auto connor_copy = get_copy(connor_ptr);
                        ^
note: constraints not satisfied
    T get_copy(T* pointer) {
      ~~~~~
note: 'std::is_copy_constructible::value' evaluated to false

```

Listing 6-28: Program using the `get_copy` template in Listing 6-27. This code doesn't compile.

The definition of `get_copy` ❶ is followed by a `Highlander` class definition, which contains a default constructor ❷ and a deleted copy constructor ❸. Within `main`, you've initialized a `Highlander` ❹, taken its reference ❺, and attempted to instantiate `get_copy` with the result ❻. Because there can be only one `Highlander` (it's not copyable), Listing 6-28 produces an exquisitely clear error message.

static_assert: The Preconcepts Stopgap

As of C++17, concepts aren't part of the standard, so they're not guaranteed to be available across compilers. There is a stopgap you can apply in the interim: the `static_assert` expression. These assertions evaluate at compile time. If an assertion fails, the compiler will issue an error and optionally provide a diagnostic message. A `static_assert` has the following form:

```
static_assert(boolean-expression, optional-message);
```

In the absence of concepts, you can include one or more `static_assert` expressions in the bodies of templates to assist users in diagnosing usage errors.

Suppose you want to improve the error messages of `mean` without leaning on concepts. You can use type traits in combination with `static_assert` to achieve a similar result, as demonstrated in Listing 6-29.

```
#include <type_traits>

template <typename T>
T mean(T* values, size_t length) {
    static_assert(std::is_default_constructible<T>(),
        "Type must be default constructible."); ❶
    static_assert(std::is_copy_constructible<T>(),
        "Type must be copy constructible."); ❷
    static_assert(std::is_arithmetic<T>(),
        "Type must support addition and division."); ❸
    static_assert(std::is_constructible<T, size_t>(),
        "Type must be constructible from size_t."); ❹
    --snip--
}
```

Listing 6-29: Using `static_assert` expressions to improve compile time errors in `mean` in Listing 6-10.

You see the familiar type traits for checking that `T` is default ❶ and copy constructible ❷, and you provide error methods to help users diagnose issues with template instantiation. You use `is_arithmetic` ❸, which evaluates to true if the type parameter supports arithmetic operations (+, -, /, and *), and `is_constructible` ❹, which determines whether you can construct a `T` from a `size_t`.

Using `static_assert` as a proxy for concepts is a hack, but it's widely used. Using type traits, you can limp along until concepts are included in the standard. You'll often see `static_assert` if you use modern third-party libraries; if you're writing code for others (including future you), consider using `static_assert` and type traits.

Compilers, and often programmers, don't read documentation. By baking requirements directly into the code, you can avoid stale documentation. In the absence of concepts, `static_assert` is a fine stopgap.

Non-Type Template Parameters

A template parameter declared with the `typename` (or `class`) keyword is called a *type template parameter*, which is a stand-in for some yet-to-be-specified type. Alternatively, you can use *non-type template parameters*, which are stand-ins for some yet-to-be-specified value. Non-type template parameters can be any of the following:

- An integral type
- An lvalue reference type
- A pointer (or pointer-to-member) type

- A `std::nullptr_t` (the type of `nullptr`)
- An `enum class`

Using a non-type template parameter allows you to inject a value into the generic code at compile time. For example, you can construct a template function called `get` that checks for out-of-bounds array access at compile time by taking the index you want to access as a non-type template parameter.

Recall from Chapter 3 that if you pass an array to a function, it decays into a pointer. You can instead pass an array reference with a particularly off-putting syntax:

```
element-type(¶m-name)[array-length]
```

For example, Listing 6-30 contains a `get` function that makes a first attempt at performing bounds-checked array access.

```
#include <stdexcept>

int& get(int (&arr)[10]❶, size_t index❷) {
    if (index >= 10) throw std::out_of_range{ "Out of bounds" };❸
    return arr[index];❹
}
```

Listing 6-30: A function for accessing array elements with bounds checking

The `get` function accepts a reference to an `int` array of length 10 ❶ and an index to extract ❷. If `index` is out of bounds, it throws an `out_of_range` exception ❸; otherwise, it returns a reference to the corresponding element ❹.

You can improve Listing 6-30 in three ways, which are all enabled by non-type template parameters genericizing the values out of `get`.

First, you can relax the requirement that `arr` refer to an `int` array by making `get` a template function, as in Listing 6-31.

```
#include <stdexcept>

template <typename T❶>
T&❷ get(T❸ (&arr)[10], size_t index) {
    if (index >= 10) throw std::out_of_range{ "Out of bounds" };
    return arr[index];
}
```

Listing 6-31: A refactor of Listing 6-30 to accept an array of a generic type

As you've done throughout this chapter, you've genericized the function by replacing a concrete type (here, `int`) with a template parameter ❶❷❸.

Second, you can relax the requirement that `arr` refer to an array of length 10 by introducing a non-type template parameter `length`. Listing 6-32 shows how: simply declare a `size_t length` template parameter and use it in place of 10.

```
#include <stdexcept>

template <typename T, size_t Length❶>
T& get (T(&arr)[Length❷], size_t index) {
    if (index >= Length❸) throw std::out_of_range{ "Out of bounds" };
    return arr[index];
}
```

Listing 6-32: A refactor of Listing 6-31 to accept an array of a generic length

The idea is the same: rather than replacing a specific type (int), you've replaced a specific integral value (10) ❶❷❸. Now you can use the function with arrays of any size.

Third, you can perform compile time bounds checking by taking size_t index as another non-type template parameter. This allows you to replace the std::out_of_range with a static_assert, as in Listing 6-33.

```
#include <cstdio>

template <size_t Index❶, typename T, size_t Length>
T& get(T (&arr)[Length]) {
    static_assert(Index < Length, "Out-of-bounds access"); ❷
    return arr[Index❸];
}

int main() {
    int fib[] { 1, 1, 2, 0 }; ❹
    printf("%d %d %d ", get<0>(fib), get<1>(fib), get<2>(fib)); ❺
    get<3>(fib) = get<1>(fib) + get<2>(fib); ❻
    printf("%d", get<3>(fib)); ❼
    //printf("%d", get<4>(fib)); ❽
}
```

```
1 1 2 ❺❸ ❼
```

Listing 6-33: A program using compile time bounds-checked array accesses

You've moved the size_t index parameter into a non-type template parameter ❶ and updated the array access with the correct name Index ❸. Because Index is now a compile time constant, you also replace the logic _error with a static_assert, which prints the friendly message Out-of-bounds access whenever you accidentally try to access an out-of-bounds element ❷.

Listing 6-33 also contains example usage of get in main. You've first declared an int array fib of length 4 ❹. You then print the first three elements of the array using get ❺, set the fourth element ❻, and print it ❼. If you uncomment the out-of-bounds access ❽, the compiler will generate an error thanks to the static_assert.